

ZHANG YILIN

(+86) 135 2045 8206 | zhang_yilin@ruc.edu.cn | github.com/zhangzzzgog | zhangzzzgog.github.io

EDUCATION

Renmin University of China (RUC) <i>Ph.D. Student, School of Information Advisors: Yueguo Chen, Chunyu Wei AI safety & applications</i> <i>Beijing, China</i>	09/2025 — Present
University of Cambridge <i>Summer Exchange Programme Neuroscience & Machine Learning</i>	07/2023 — 08/2023 Cambridge, UK
Wuhan University (WHU) <i>B.Eng., Computer Science and Technology HongYi Honor School LLM security</i>	09/2021 — 06/2025 Wuhan, China

PUBLICATIONS

Papillon: Efficient and Stealthy Fuzz Testing-Powered Jailbreaks for LLMs <i>USENIX Security 2025 — Published</i>	07/2023 — 09/2024
To overcome the poor adaptivity and scalability of existing LLM jailbreak attacks, proposed a two-stage, fuzz-testing-based self-mutation method for adaptive jailbreaking that attacks efficiently with zero dependence on seed templates.	
Don't Click That: Teaching Web Agents to Resist Deceptive Interfaces <i>ACL 2026 — Published (First Author)</i>	11/2025 — 01/2026
Built a benchmark of complex web scenes mixing benign and deceptive elements, and designed an RL scheme driven by the agent's click confidence with an experience pool for consolidation. The resulting plug-and-play visual module reduces susceptibility to deceptive interfaces without making agents over-conservative.	

RESEARCH IN PROGRESS

Rethinking the Role of Retrieval in Agent Memory <i>Under review (First Author)</i>	02/2026 — 05/2026
Found that the key to agent-memory retrieval lies in the sparse expression of a few salient memory elements, a signal that dense vector retrieval flattens. Proposed a retrieval paradigm incorporating sparse-element statistics, outperforming complex memory architectures in both recall and generation quality.	
Hierarchical Disclosure of Sensitive Knowledge Inside LLMs <i>Under review (First Author)</i>	03/2026 — 07/2026
Sensitive training data in domain LLMs is disclosed regardless of whether the asker is authorized. Designed a tiered-disclosure evaluation set and a probe-guided control module that makes the model follow authorization-aware disclosure with very few parameter updates.	
Post-hoc Watermarking via Localized Re-diffusion of Existing Text <i>Under review (First Author)</i>	08/2026 — 09/2026
Recast post-hoc watermarking of already-generated text as the unmasking process of a masked diffusion model, and designed a localized re-diffusion paradigm that achieves efficient and stable watermark embedding and detection.	

PROJECT EXPERIENCE

LLM-Based Financial Analysis & Stock Trading System <i>Group Project (Lead) Python · Vue.js · LSTM</i>	09/2024 — 11/2024
Led the design of an end-to-end stock-analysis platform combining a real-time news crawler, LLM sentiment analysis and LSTM price prediction, with a Vue.js dashboard for signal monitoring and backtesting.	
Multinational Corporate Investment Simulation via Multi-Agent System <i>Group Project (Lead) Python · ReAct</i>	11/2025
Two-layer multi-agent system: a dispatch agent decomposes queries into company-level subtasks; company agents simulate annual FDI decisions with personalized parameters. Simulated trends across 62 companies, validated against real FDI data.	

SKILLS & OTHERS

Skills: Python, C++; Claude Code, Codex, Visual Studio Code
Languages: Mandarin (Native), English (TOEFL 103, GRE 320; Fluent)
Awards: Third Prize, Chinese National Chemistry Olympiad; First Prize, Beijing Chemistry Olympiad